



# Performance of Docker vs VMs

---

Presented By  
Ali Hussain

August 21, 2014

# Today's Presenter:

---

## Ali Hussain

Co-founder & CTO Flux7

Prev: CPU Performance Analyst at Intel and ARM

### Flux7: Cloud and DevOps Solutions

Cloud and Devops for Web teams

Enterprise DevOps management

AWS Certified Team



---

#### Partners:



#### Clients:



# Other team members

---

## Samprita Hegde

Performance Engineer at Flux7

Setup, execution, and collection of data

Check out our work on [blog.flux7.com](https://blog.flux7.com)



# Agenda

---

- —
- —
- —

Background

- —
- —
- —

Experimental results & deductions


- —
- —
- —

Application in the real world

# What is Docker?

---

Linux OS isolation  
tools made easy



A Docker container  
looks like a  
virtual machine



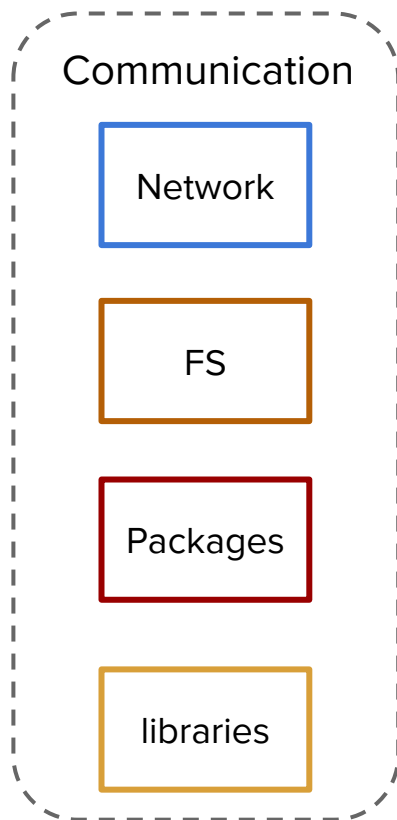
Provide additional  
'goodies' for app  
development

# Holy Grail of “Virtualization”

Host  
Hardware  
and OS

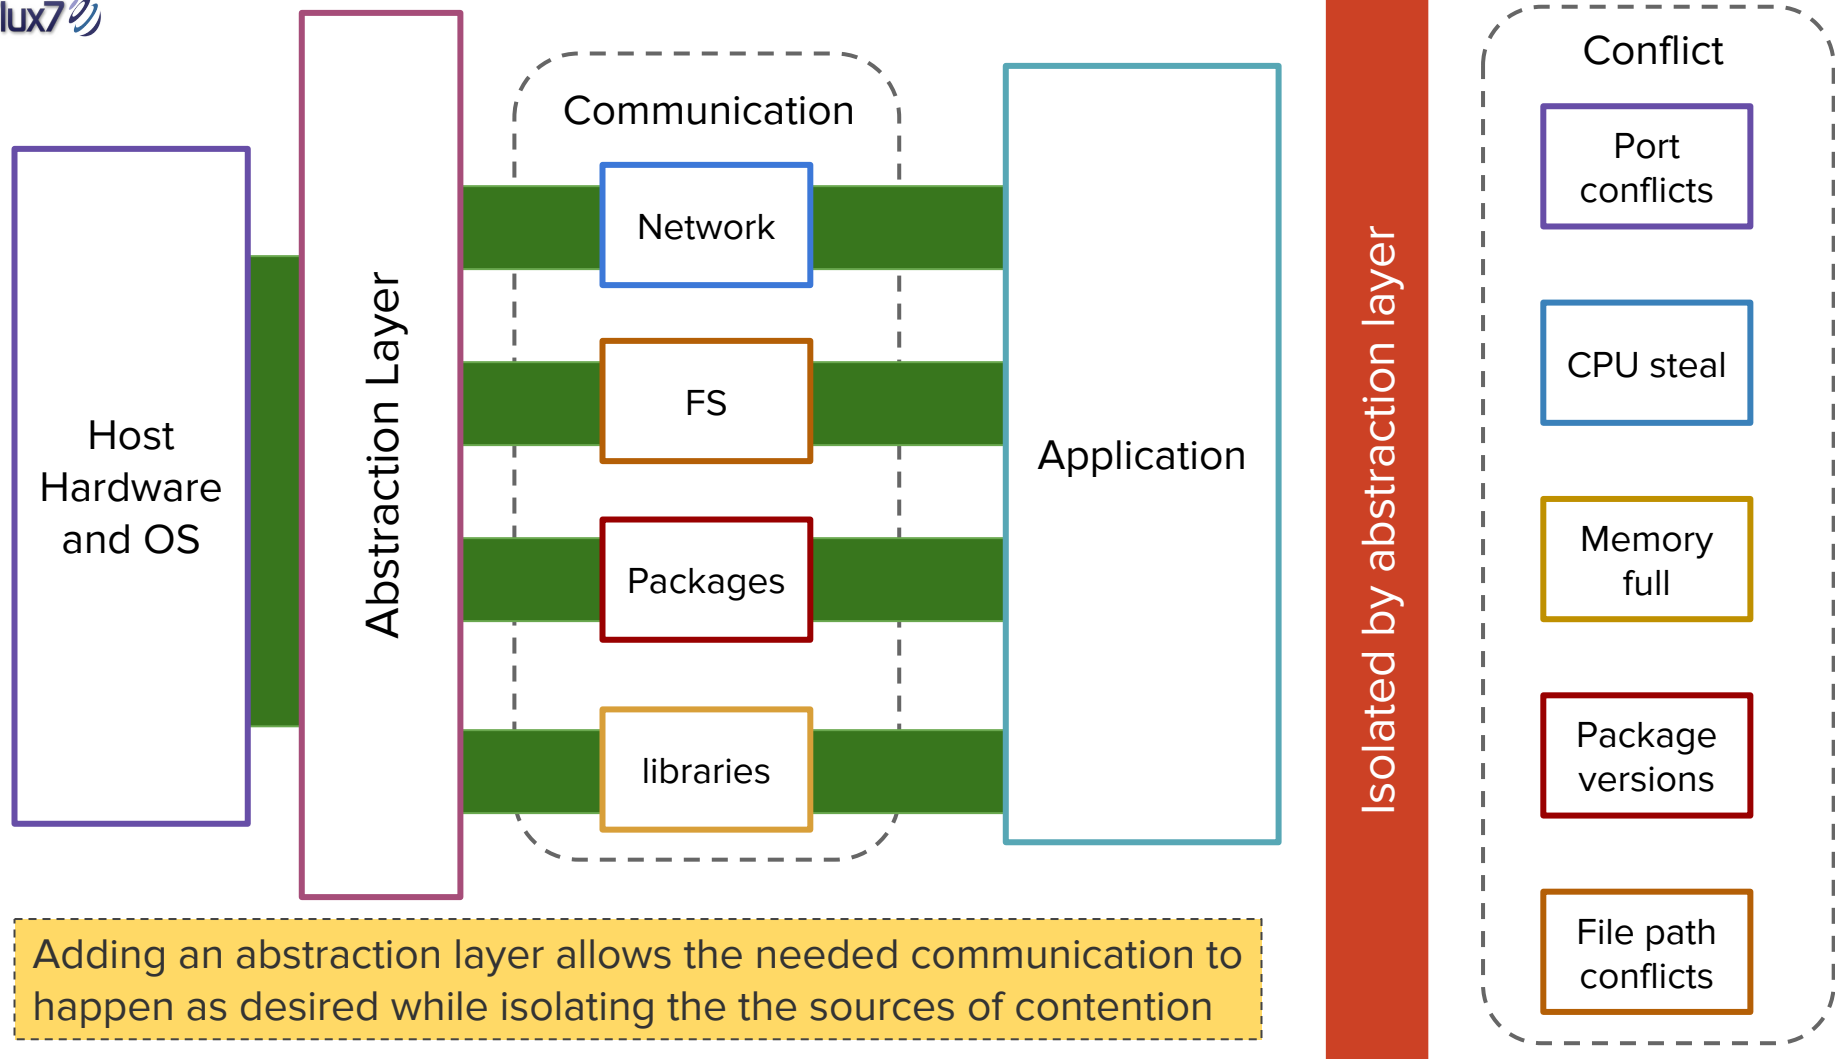
Application

An application needs to communicate with  
outside world: H/W, OS, other processes



The application communicates through certain channels but each of these channels can also cause conflicts with other applications on the system





# What we want

---



Allow  
communication




Prevent conflict  
over resources



A continuum of  
abstraction levels

# How we pay for it

---




Resource overhead  
(Memory, CPU)



Increased latency  
(Disk, Network)



Snapshot time and  
space




Start and stop time

# Where is Docker on this spectrum?

---

Virtual  
Machines



Native

# Docker Abstraction

---

- ✓ Use isolation features of the Linux kernel to give a VM-like environment
- ✓ Filesystem using chroot jails
- ✓ Network using port forwarding
- ✓ Resource and process isolation using cgroups

# Test Platform



**OS:** Ubuntu 12.04 LTS , kernel - 3.8.0-33-generic



**Docker version:** 0.7.2, build 28b162e

## Processor

Intel(R) Core  
(TM) i7 CPU

## Memory

47 GB RAM  
23 GB Swap

## Disk

2 drive raid1  
(mirroring)  
configuration

\* Some tests run on AWS because of specific needs

# Start and stop times

---

	Start Time	Stop Time
<b>Docker Containers</b>	<50ms	<50ms
<b>VMs</b>	30-45 seconds	5-10 seconds

# Memory

---



No overhead of running a guest kernel



Docker requires consumed memory not provisioned memory



# CPU Overhead - Methodology

---




Measured using coremark



Microbenchmark measuring CPU performance

# CPU Overhead

---



# CPU Overhead

---



No noticeable difference between native, KVM, and Docker performance



KVM uses ~1.5% more CPU usage compared to Docker when idle

# Network Performance - Methodology

---




Performance measured using iperf




Care about both latency and throughput

# Network Throughput



**Theoretical  
Scenario:**  
Server and Client  
on same instance

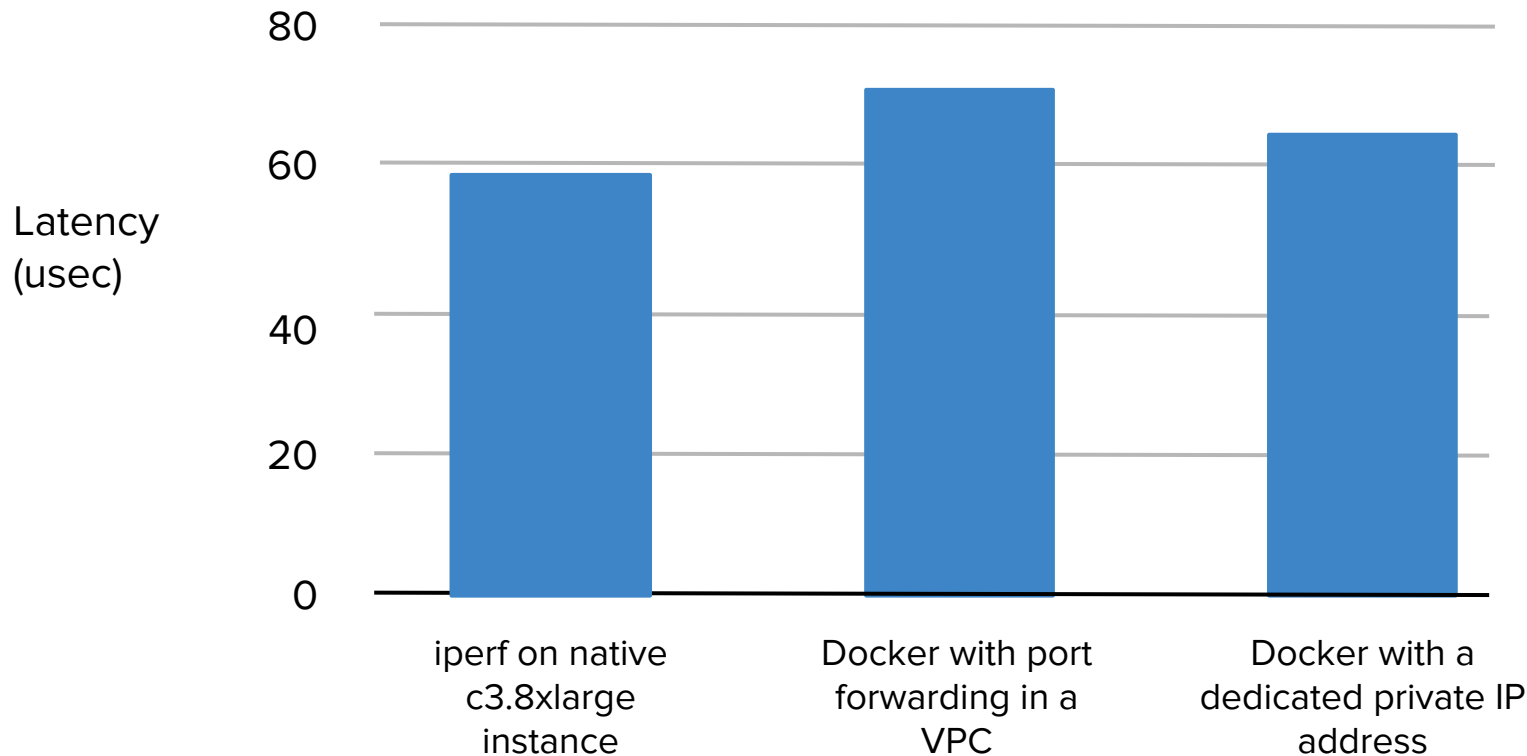
# Network Throughput



**Practical Scenario:**  
Server and Client  
on different  
instances

# Network latency (lower is better)

---



# Network Performance

---



Both Docker and KVM can saturate 10GbE



Docker running inside a VM can still saturate 10GbE



Docker redirection latency  $\sim 10\mu\text{s}$




# Disk Bandwidth

---



Measured using FIO - Stresses disk with different streams

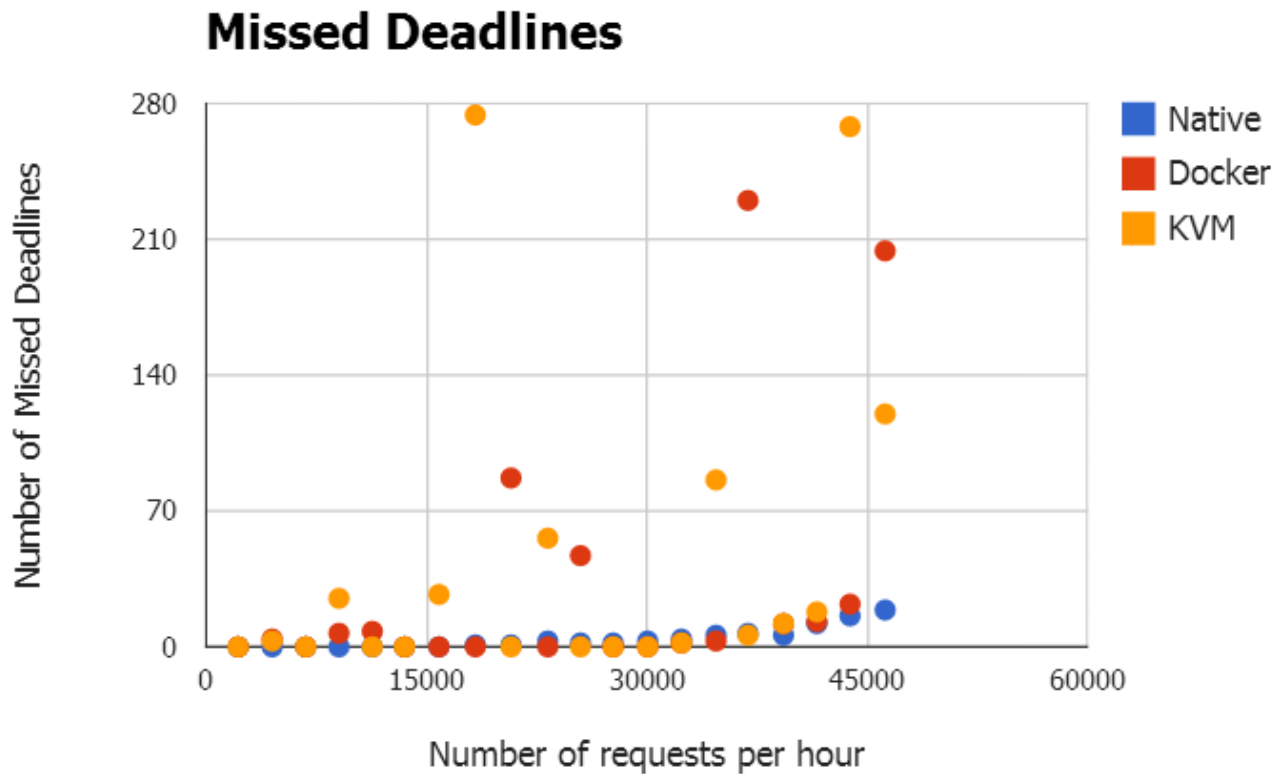


A lot of variation but no clear winner



Different drivers causing differences with different scenarios

# Application Benchmarking - Wikibench



# Benchmark summary


---

Load performance:

- ✓ Native, KVM, and Docker performed comparably
- ✓ H/W and S/W advances in virtualization
- ✓ Over micro and macro benchmarks

Docker shines in

- ✓ Idle resource usage
- ✓ Start and stop times




# The nature of disruptive changes

New technologies provide benefit only if  
you drop the rules that were designed to  
overcome the shortcomings of the old  
technology

-- Dr. Eliyahu Goldratt

# Use Cases

---



Configuration  
Management



Multiple instances  
per box



Basic unit of  
provisioning  
resources

# Configuration Management

---



Create configuration to work in your canonical environment and emulate it



We already do this

- Virtual machines
- Python virtualenvs

# Configuration Management

---

## Why Docker?



Consistent  
environment



Many dev niceties



No memory overhead



Minimize performance  
overhead

# Configuration Management

---

What we get?



Running identical code + configs across environments



Better code pipeline management



# Use Cases

---



Configuration  
Management



Multiple instances  
per box



Basic unit of  
provisioning  
resources

# Multiple Instances per box

---



Multiple instances each running in own Docker container



Containers may communicate over TCP

# Multiple Instances per box

---

## Why Docker?



Process and FS  
isolation



Port forwarding



No idle memory and  
CPU overhead

# Multiple Instances per box

---

What we get



High fidelity local  
dev environments



Multi-tenancy



Server consolidation

# Use Cases

---



Configuration  
Management



Multiple instances  
per box



Basic unit of  
provisioning  
resources

# Basic Unit of Resource Provisioning

---



Allocate Docker containers instead of VMs

Can be both



Persistent resources



Very short tasks

# Basic Unit of Resource Provisioning

---

## Why Docker



Sharing of  
resources



Quick start and  
stop times



Management of  
multiple images

# Basic Unit of Resource Provisioning

---

What we get



Faster provisioning



Lower overhead on  
host



Config management



# Thank you, Docker!

For developing a very cool piece of technology

For making it OpenSource

Special thanks to **Jerome Petazzoni**  
for supporting us





# Thank You!



**Twitter:** [@Flux7Labs](https://twitter.com/Flux7Labs)



**Blog:** [blog.flux7.com](https://blog.flux7.com)